

# Report 2: Pseudo Random Generation of Numbers

Generating random numbers is very important for cryptography. So, a computer who utilizes these techniques very often must have a robust and secure way of generating them. In this report we will explore a myriad of techniques and see how secure an entropic they are.

## Task 1: Generating an Encryption Key the Wrong Way

Firstly, we will generate an **Encryption key** the wrong way, and for that, we will use a program that uses the random library of c and the time to generate random numbers. The following is the program used:

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

#define KEYSIZE 16

void main()
{
    int i;
    char key[KEYSIZE];
    printf("%lld\n", (long long) time(NULL));
    srand (time(NULL));
    for (i = 0; i < KEYSIZE; i++){
        key[i] = rand()%256;
        printf("%.2x", (unsigned char)key[i]);
    }
    printf("\n");
}
```

The program generates a 128 bit char random **Encryption key** using srand and a seed.

As we can see, we generate the seed for the random function in the line **srand (time(NULL));**. We use the current time as the seed, and then in the line **key[i] = rand()%256;** we generate a number between 0 and 256 to add to the **Encryption key**. **time(NULL)** returns the time in seconds elapsed from January 1, 1970 to now. By running the program we generate the following keys:

```
[03/05/24] seed@VM:~/.../Lab3$ ./task1
1709638435
7c319c868cf47f8429dab249acbf153f
[03/05/24] seed@VM:~/.../Lab3$ ./task1
1709638504
709c3bdc349e3a4f7ea991cde3e0ce05
[03/05/24] seed@VM:~/.../Lab3$
```

If we remove the line defining the seed as the current time, the program uses the default seed for every execution, which is **1**. So the resulting **Encryption keys** are always the same, as we can see in the image below:

```
[03/05/24] seed@VM: ~/.../Lab3$ ./task1
1709638173
67c6697351ff4aec29cdbaabf2fbe346
[03/05/24] seed@VM: ~/.../Lab3$ ./task1
1709638175
67c6697351ff4aec29cdbaabf2fbe346
[03/05/24] seed@VM: ~/.../Lab3$ ./task1
1709638176
67c6697351ff4aec29cdbaabf2fbe346
[03/05/24] seed@VM: ~/.../Lab3$ █
```

## Task 2: Guessing the Key

So, what is the problem with this way of doing it this way? Let us say that Alice generated an **Encryption key** using the program above between the date of **2018-04-17 21:08:49** and **2018-04-17 23:08:49** and used it to encrypt some very important documents. As we saw before, the program above uses time as a seed to generate the **Encryption key**, so, as we know the timeframe of the generation we can execute the program and obtain the same seed as Alice! We can then obtain her **Encryption key** and decrypt all her important files! We will now attempt to get Alice's documents. Let us assume that the documents are in the **pdf** format. The document is encrypted using **AES**, which is a 128-bit cypher that uses blocks to encrypt and decrypt documents. A block consists of 16 bytes and so we will need 16 bytes of plaintext to decrypt it, so how do we get them? But first we need to get Alice's **Encryption Key**, and so we need to generate every possible key in that 2 hour period and test against some plain text we can find in the **pdf** file. We are lucky as we can get 16 bytes of the header of the files. The beginning part of a **pdf** header is always the version number. At the time the file was created **PDF-1.5** was widely used so we now have 8 bytes of clear data in the form of the version number **%PDF-1.5**. The next 8 bytes are also quite easy to predict and so we now have 16 bytes of plain text. We compare the result to the encrypted data of those 16 bytes and see if it matches. The **plaintext** and **ciphertext** are below:

```
Plaintext: 255044462d312e350a25d0d4c5d80a34
Ciphertext: d06bf9d0dab8e8ef880660d2af65aa82
IV: 09080706050403020100A2B2C2D2E2F2
```

We also have access to the **Initial Vector(IV)** as it is never encrypted. With all that done we can finally start the brute force attack using the following programs:

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
```

```

#define KEYSIZE 16

void main()
{
    int i;
    FILE* f1;
    f1 = fopen("enckey.txt", "a+");

    for (i = 1524013729; i <= 1524020929; i++) {
        char key[KEYSIZE];
        printf("%lld\n", (long long) i);
        srand(i);

        int j;

        for (j = 0; j < KEYSIZE; j++) {
            key[j] = rand()%256;
            printf("%.2x", (unsigned char) key[j]);
            fprintf(f1, "%.2x", (unsigned char) key[j]);
        }
        fprintf(f1, "\n");
        printf("\n");
    }
}

```

This c program computes from every possible seed (the 2 hour window) starting in **1524013729** and ending in **1524020929**, computes every **Encryption Key** and saves it to the file **enckey.txt**. Now with all the possible keys we can check which one is actually the right one. And for this we use the python script bellow:

```

from Crypto.Cipher import AES

file = open("enckey.txt", "r")

Lines = file.readlines()

for key in Lines:
    aes = AES.new(bytearray.fromhex(key), AES.MODE_CBC,
bytearray.fromhex("09080706050403020100A2B2C2D2E2F2"))
    data = aes.encrypt(bytearray.fromhex("255044462d312e350a25d0d4c5d80a34"))
    if data == bytearray.fromhex("d06bf9d0dab8e8ef880660d2af65aa82"):
        print("Key gotten!\n")
        print(key)

```

This python scrip imports the **AES** algorithm, gets all the keys from the **enckey.txt** and creates an **AES** with each key and with the **IV** collected, it then encrypts our 16 bytes of plaintext and compares it with the equivalent cyphertext. If they are the same that means we found the actual key used in the encryption! We can see the python script execute in the image bellow:

```
[03/05/24] seed@VM:~/.../Lab3$ python3 cracker.py  
Key gotten!
```

```
95fa2030e73ed3f8da761b4eb805dfd7
```

So the Encryption Key is **95fa2030e73ed3f8da761b4eb805dfd7**! We found Alice's key and can now decrypt all her encrypted files.

### Task 3: Measure the Entropy of Kernel

It is difficult for computers to create randomness. So most Operating Systems get their randomness from the physical world. Linux gets the randomness from these physical resources:

```
void add_keyboard_randomness(unsigned char scancode);  
void add_mouse_randomness(__u32 mouse_data);  
void add_interrupt_randomness(int irq);  
void add_blkdev_randomness(int major);
```

The **OS** is using the timing between keypresses to generate random numbers, the movement of the mouse and interrupt timing, the interrupt timing of the disk and finally the finish time of block device requests. But how can we judge the quality of the randomness? We can do it with **Entropy**, it measures how many bits of random numbers the system has. In Linux you can check the system's entropy in the file **/proc/sys/kernel/random/entropy\_avail**. We can check it by doing the following command:

```
$ cat /proc/sys/kernel/random/entropy_avail
```

This command returns the number of bytes of entropy available for random number generation, we can see how many we have in the image bellow:

```
[03/05/24] seed@VM:~/.../Lab3$ cat /proc/sys/kernel/random/entropy_avail  
1523  
[03/05/24] seed@VM:~/.../Lab3$ cat /proc/sys/kernel/random/entropy_avail  
1534  
[03/05/24] seed@VM:~/.../Lab3$ cat /proc/sys/kernel/random/entropy_avail  
1543  
[03/05/24] seed@VM:~/.../Lab3$  
[03/05/24] seed@VM:~/.../Lab3$ cat /proc/sys/kernel/random/entropy_avail  
1586  
[03/05/24] seed@VM:~/.../Lab3$ █
```

As we can see, we have around 1500 bytes of entropy available. But, as stated before, the levels of entropy can be raised by using the keyboard, mouse and other methods. By executing the following command we can check how the number of bytes change by using the keyboard and mouse.

```
$ watch -n .1 cat /proc/sys/kernel/random/entropy_avail
```

The command checks the changes in the **/proc/sys/kernel/random/entropy\_avail** file every 0.1 seconds. We can see in the image below that after we move the mouse and press some keys the entropy value increases.

```
Every 0.1s: cat /proc/sys/kernel/random/entropy... VM: Tue Mar  5 12:48:59 2024
2317
```

Through this experiment we can see how the system obtains a reliable source of entropy to generate resilient pseudo random numbers.

## Task 4: Get Pseudo Random Numbers From /dev/random

Linux uses two devices to use the physical numbers collected by **/proc/sys/kernel/random/entropy\_avail**. Those two being **/dev/random** and **/dev/urandom**. They behave differently, as **/dev/random** blocks and does not generate any number when the **/proc/sys/kernel/random/entropy\_avail** has 0 bytes of collected data. **/dev/random** will only resume operation when it finds that the physical data is enough for generating pseudo random numbers. We can see the behaviour of the **/dev/random/** by executing the command:

```
$ cat /dev/random | hexdump
```

**Hexdump** makes the output more understandable. When we run the output we see the following:

```
[03/05/24]seed@VM:~$ cat /dev/random | hexdump
00000000 4de8 6d5f a70b 2ba7 b14c 6883 6d89 4cca
00000010 f139 6ad8 3b54 f312 eb12 767b 9517 f628
00000020 504d 6b92 9e90 e0b5 7067 35da 3839 3751
00000030 2880 9d35 2db8 2a1c f9c3 a571 46b9 05f6
00000040 f2b7 87d7 529d e908 fe19 374f da66 b3ae
00000050 7c6d 2982 4f11 8ca2 1170 6b99 bd2e f185
00000060 2d19 d773 c046 dfcd 6db7 7b39 3e5c 5b0c
```

When we execute the watch on **cat /proc/sys/kernel/random/entropy\_avail** and see the entropy increases, we also see that more lines are added to **/dev/random** until it goes to 0. When enough bytes are collected another entry is made on **/dev/random** and the entropy goes back to 0 again. We can see the blocking behaviour here, as entries are only made when there are enough bytes collected. **/dev/random** should be used in limited amounts because of this blocking behaviour. For example, if we use **/dev/random** in a server to generate random session keys we can cause a **DOS**. If the number of clients is greater than the number of lines in **/dev/random** has and can produce, the server will block waiting for the entropy and the clients are left waiting. This can then constitute a **DOS**. A malicious attacker can also send many false client requests and can bring the server down to legitimate clients.

## Task 5: Get Pseudo Random Numbers From /dev/urandom

Linux also provides another device for pseudo random numbers and that is the **/dev/urandom**. The difference to **/dev/random** is that it will not block when the entropy in **/proc/sys/kernel/random/entropy\_avail** is 0. **/dev/urandom** uses the data pooled to create a seed for its random number generation. Let us see the behaviour of **/dev/urandom** by, once again, executing the following command:

```
$ cat /dev/urandom | hexdump
```



---

030ad90	73c2	6e32	2c5e	2e23	fb37	c204	ebbb	f5e7
030ada0	ff77	b192	def6	8be0	8ab3	34fa	10e9	c0ff
030adb0	2064	9b03	dc2d	a4a9	94e8	409b	e495	8f54
030adc0	78f9	b3f9	38ef	83af	ef37	a400	c4e8	55f1
030add0	634d	cffd	39bb	37e2	2c49	dc1d	c244	2c34
030ade0	4b38	1355	6b00	cd2b	c17f	cfbf	d246	3753
030adf0	37bd	55b6	2111	674f	9494	1614	a4d1	7fd8
030ae00	30c4	32ac	c160	1721	866c	7c6b	406b	3ebe
030ae10	c5fe	1a2b	cb70	f743	8bc3	13cb	5fcb	72be
030ae20	dd57	f885	b5e8	a4de	4cb8	f325	05c9	6655
030ae30	37af	e46a	1fa6	5d15	cfbe	1109	fd89	15f3
030ae40	8692	b1cf	a2b1	c749	f253	5a5a	7b4e	454d
030ae50	e14c	c54c	ff0f	761d	45c0	5df9	09f2	27e1
030ae60	92e6	c3a9	5f58	6189	f7b5	57f3	70a6	28fa
030ae70	b519	530b	8365	21c4	c83c	ed77	712f	44a0
030ae80	9033	b5bd	321b	a1a1	8d33	477a	58dc	1ffd
030ae90	47a0	c05a	d1ed	189e	0b4d	f0aa	52bb	b1cc
030aea0	3072	f3c3	70bd	af3e	c839	f424	9bab	d7e6
030aeb0	7d59	bc40	e8d8	4436	e039	2ba3	3d30	dd59
030aec0	e598	4b6e	5479	b798	a979	b2a0	4bac	0822
030aed0	067c	a5d8	d21f	f6fe	4248	8622	2c5a	69b3
030aee0	438a	8219	7e22	3445	f29b	b925	1633	c5a3
030aef0	7d05	a235	e6e2	a259	7064	bbb1	82d2	490b

Immediately we see that it generates numbers non-stop, as it is non blocking, when the entropy runs out it starts using the pool instead. But are the results any good? We can check their quality using a tool called **ENT**, it analyses the numbers generated in a myriad of tests. Let's test 1 MB of the pseudo numbers generated by **/dev/urandom**. To test the 1 MB we will use the following commands:

```
$ head -c 1M /dev/urandom > output.bin
$ ent output.bin
```

The first command generates 1 MB of data and saves it to **output.bin**. The second tests it using **ENT**. The following image shows the results:

```
[03/05/24]seed@VM:~$ head -c 1M /dev/urandom > output.bin
[03/05/24]seed@VM:~$ ent output.bin
Entropy = 7.999832 bits per byte.
```

Optimum compression would reduce the size of this 1048576 byte file by 0 percent.

Chi square distribution for 1048576 samples is 244.91, and randomly would exceed this value 66.39 percent of the times.

```
Arithmetic mean value of data bytes is 127.4410 (127.5 = random).
Monte Carlo value for Pi is 3.153156865 (error 0.37 percent).
Serial correlation coefficient is -0.001467 (totally uncorrelated = 0.0).
[03/05/24]seed@VM:~$ █
```

The test shows that the results are quite good! They are not totally random but are very close as we can see in the **arithmetic mean value**, **monte carlo** and **serial correlation coefficient**. They are close to the best value possible. In theory, **/dev/random** is more secure as it gets the best results but, for practical applications, we use **/dev/urandom** as the blocking penalty is very costly because it causes **Denial of Service**. So, using what we learned, let's modify the program that generates **Encryption Keys** using **/dev/urandom**. The following is the modified code:

```
#include <stdio.h>
#include <stdlib.h>

#define KEYSIZE 32

void main()
{
    int i;
    unsigned char* key = (unsigned char *) malloc(sizeof(unsigned char) *
KEYSIZE);
    FILE* random = fopen("/dev/urandom", "r");

    fread(key, sizeof(unsigned char) * KEYSIZE, 1, random);
    fclose(random);

    for (i = 0; i < KEYSIZE; i++) {
        printf(".2x", (unsigned char) key[i]);
    }

    printf("\n");
}
```

The script uses **/dev/urandom** to generate a 256-bit Encryption Key. It accesses the device and reads 32 bytes of data. This way the **Encryption key** is resilient to bruteforce attacks of the type used before due to using a pseudo random number and not a predictable seed. We can see the key generated in the following image:



```
[03/05/24] seed@VM:~/.../Lab3$ sudo ./task1
754b7054d640d15da701e8ac250ee771d4626e40d007b88b7c9c685845f3c27b
[03/05/24] seed@VM:~/.../Lab3$ sudo ./task1
0c2c407a819d6a9a6158579db869d57199d1942b667dfce40130ad3a2ea3358c
```

Authored by:

Eduardo Ramos, up201906732  
Leandro Silva, up202008061  
Luís Paiva, up202006094